

Under a strong (ToG-2.0-style) engine, the LLM judge wins in 3 of 4 regimes; embedding wins only at 1-hop

